

# Introduction to FreeRTOS

## Part 2

## **Interrupt Management**

### **Events**

Embedded real-time systems have to take actions in response to events that originate from the environment.

How should they be detected? Interrupts, polling

What kind of processing needs to be done? Inside ISR, outside ISR

### **Interrupt priority vs task priority**

Lowest priority interrupt pre-empt highest priority task

### **Interrupt Safe API Function**

FreeRTOS provides two versions of some API functions:

one for use from tasks,

and one for use from ISRs (“FromISR” appended to their name).

## Interrupt Management

### Context Switching from ISR

Context switching is performed when the task that comes into running state when interrupts finish, is different from the task that went to blocked state when interrupt occurred.

Use **pxHigherPriorityTaskWoken** parameter

a variable to inform the application writer that a context switch should be performed

PdTRUE, PdFalse

### Macros to **request a context switch from ISR**

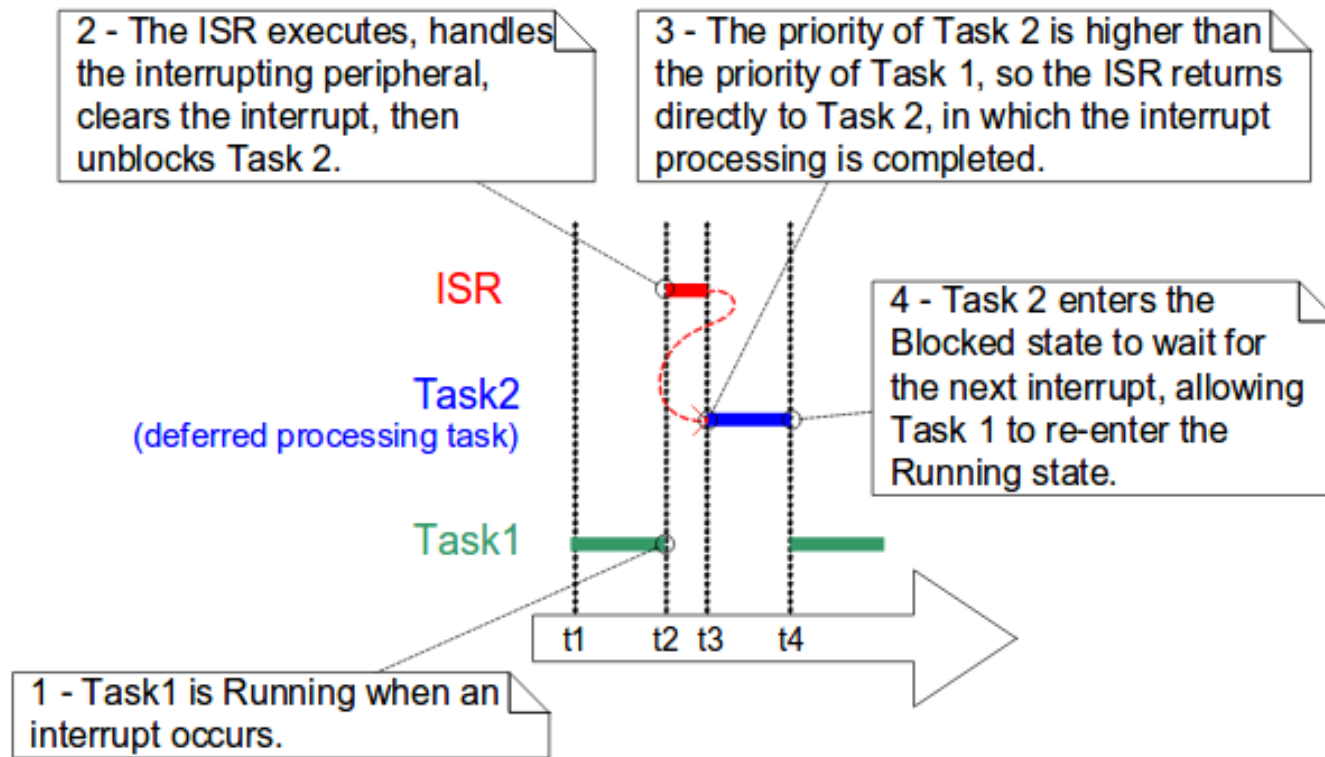
```
portYIELD_FROM_ISR( xHigherPriorityTaskWoken );
```

```
portEND_SWITCHING_ISR( xHigherPriorityTaskWoken );
```

### Macro to **request context switching from a task**

```
taskYIELD( )
```

## Interrupt Management



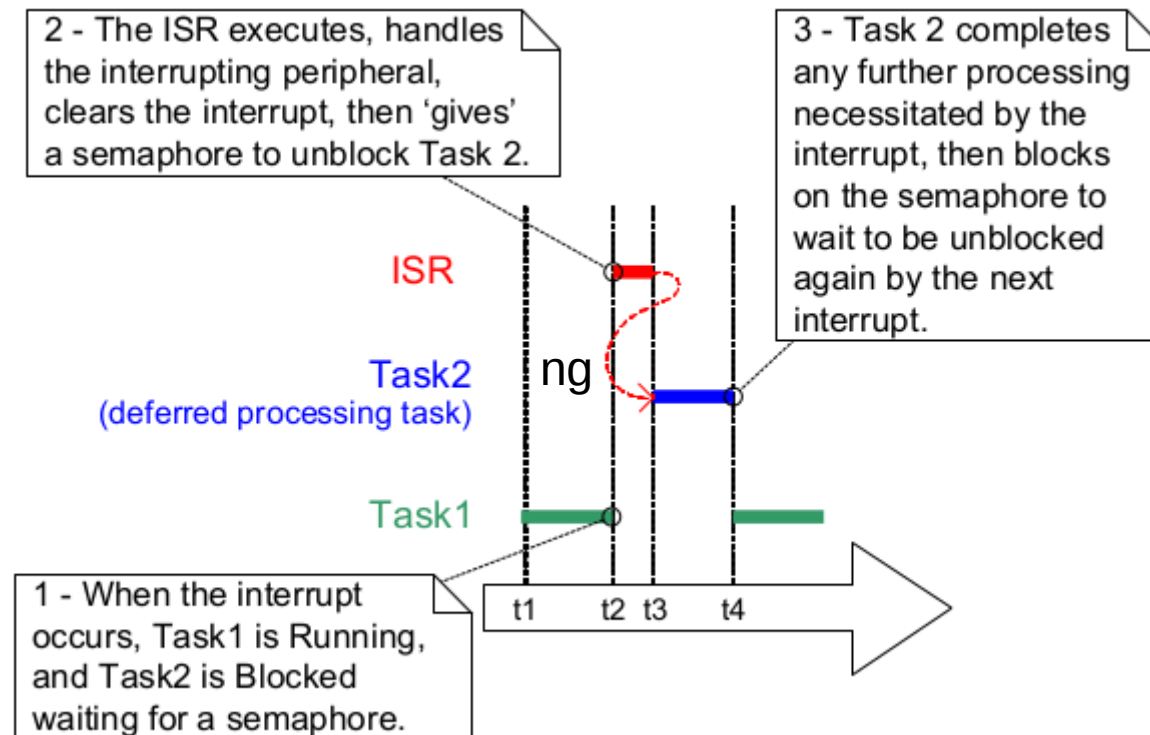
Interrupts should be deferred to a task

## Interrupt Management

### Binary Semaphores Used for Synchronization

The deferred processing task can be controlled using a ISR

- The deferred task is blocked with a “take” call for a semaphore.
- The ISR “gives” a semaphore to unblock the deferred task



## Interrupt Management

### API Functions for managing binary semaphores

#### Creating a semaphore

```
SemaphoreHandle_t  xSemaphoreCreateBinary(void);
```

#### Take

```
BaseType_t  xSemaphoreTake(SemaphoreHandle_t  xSemaphore,  
                           TickType_t  xTicksToWait);
```

#### Give

```
BaseType_t  xSemaphoreGiveFromISR(SemaphoreHandle_t  xSemaphore,  
                                   BaseType_t  *pxHigherPriorityTaskWoken);
```

## Interrupt Management

### Example Using binary semaphores (include “semphr.h”)

```
static uint32_t ulExampleInterruptHandler( void )
{
    BaseType_t xHigherPriorityTaskWoken;

    /* The xHigherPriorityTaskWoken parameter must be initialized to pdFALSE as it will get set
       to pdTRUE inside the interrupt safe API function if a context switch is required. */
    xHigherPriorityTaskWoken = pdFALSE;

    /* This interrupt does nothing more than demonstrate how to synchronise a
       task with an interrupt. A semaphore is used for this purpose. Note
       lHigherPriorityTaskWoken is initialised to pdFALSE. */
    xSemaphoreGiveFromISR( xBinarySemaphore, &xHigherPriorityTaskWoken );

    /* If there was a task that was blocked on the semaphore, and giving the
       semaphore caused the task to unblock, and the unblocked task has a priority
       higher than or equal to the currently Running task (the task that this
       interrupt interrupted), then lHigherPriorityTaskWoken will have been set to
       pdTRUE internally within xSemaphoreGiveFromISR(). Passing pdTRUE into the
       portYIELD_FROM_ISR() macro will result in a context switch being pended to
       ensure this interrupt returns directly to the unblocked, higher priority,
       task. Passing pdFALSE into portYIELD_FROM_ISR() has no effect. */
    portYIELD_FROM_ISR( xHigherPriorityTaskWoken );
}
```

## Interrupt Management

### Example Using binary semaphores

```
static void vHandlerTask( void *pvParameters )
{
    /* As per most tasks, this task is implemented within an infinite loop. */
    for( ;; )
    {
        /* Use the semaphore to wait for the event. The semaphore was created
        before the scheduler was started, so before this task ran for the first
        time. The task blocks indefinitely, meaning this function call will only
        return once the semaphore has been successfully obtained - so there is
        no need to check the value returned by xSemaphoreTake(). */
        xSemaphoreTake( xBinarySemaphore, portMAX_DELAY );
        /* To get here the event must have occurred. Process the event (in this case, just
        print out a message). */
        vPrintString( "Handler task - Processing event.\r\n" );
    }
}
```



## Interrupt Management

**Example:** Create an application that executes a periodic task reading the buttons value of the board. If one button is pressed will trigger an interrupt. This interrupt will launch a deferred task using a semaphore. The deferred task will print a message indicating which button have been pressed.

Using 4 tasks

**A periodic task reading buttons**

**An lower priority task showing a counter in leds**

**An interruptHandler task (ISR)**

**The Handlertask (deferred)**

**And the main() function**

Use the following functions (in xgpiops.h and xScuGic.h)

```
void XGpioPs_IntrEnable(XGpioPs *InstancePtr, u8 Bank, u32 Mask);
void XGpioPs_SetIntrType(XGpioPs *InstancePtr, u8 Bank, u32 IntrType,
                        u32 IntrPolarity, u32 IntrOnAny);

void XGpioPs_SetDirection(XGpioPs *InstancePtr, u8 Bank, u32 Direction);
u32 XGpioPs_Read(XGpioPs *InstancePtr, u8 Bank);
xScuGic xInterruptController;
```

## Interrupt Management

**Example:** Create an application that executes a periodic task reading the buttons value of the board. If one button is pressed will trigger an interrupt. This interrupt will launch a deferred task using a semaphore. The deferred task will print a message indicating which button have been pressed.

```
/* Scheduler includes. */  
#include "FreeRTOS.h"  
#include "task.h"  
#include "queue.h"  
#include "semphr.h"
```

```
/* Xilinx includes. */  
#include "xgpio.h"  
#include "xscugic.h"  
#include "xil_exception.h"
```

```
#include "semphr.h"
```

## Interrupt Management

Using **counting semaphores** (queues with length > 1)

It matters only the amount of items in the queue.

```
configUSE_COUNTING_SEMAPHORES = 1
```

Two functions: a) counting events  
b) resource managements

```
SemaphoreHandle_t  xSemaphoreCreateCounting( UBaseType_t  uxMaxCount,  
                                              UBaseType_t  uxInitialCount );
```

Returned value	If NULL is returned, the semaphore cannot be created because there is insufficient heap memory available
----------------	--

## Interrupt Management

### Using a queue (writing) from an interrupt

```
BaseType_t xQueueSendToFrontFromISR(QueueHandle_t xQueue,  
                                     void *pvItemToQueue,  
                                     BaseType_t *pxHigherPriorityTaskWoken );
```

```
BaseType_t xQueueSendToBackFromISR(QueueHandle_t xQueue,  
                                    void *pvItemToQueue,  
                                    BaseType_t *pxHigherPriorityTaskWoken );
```

- xQueue: The handle of the queue
- pvItemToQueue: A pointer to the data to be copied into the queue
- pxHigherPriorityTaskWoken : a variable to inform the application writer that a context switch should be performed

Return:

- pdPASS – OK
- errQUEUE\_FULL – Error, queue full

## Interrupt Management

### Using a queue (reading) from an ISR

```
BaseType_t xQueueReceiveFromISR( QueueHandle_t xQueue,  
                                void *pvBuffer,  
                                BaseType_t *pxHigherPriorityTaskWoken );
```

- xQueue: The handle of the queue
- pvBuffer: A pointer to the memory into which the data will be copied
- pxHigherPriorityTaskWoken: a variable to inform the application writer that a context switch should be performed

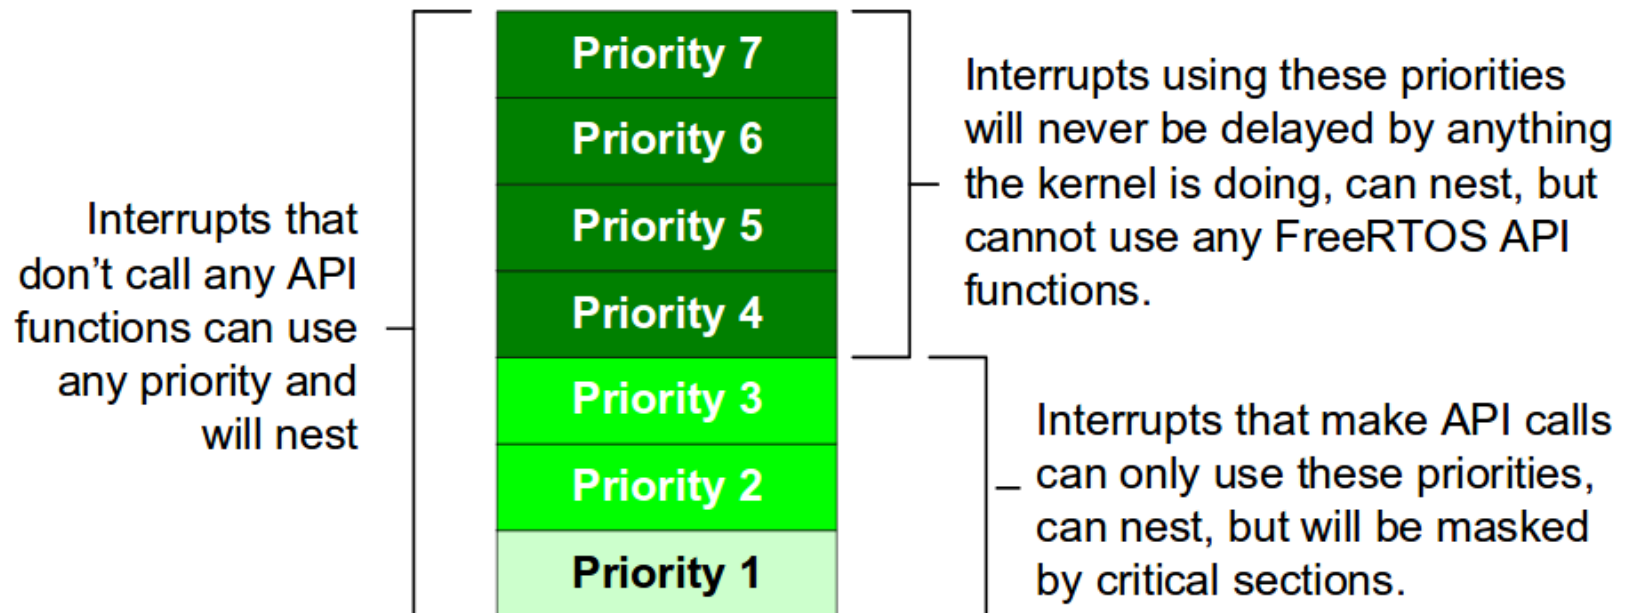
Return:

- pdPASS – OK
- errQUEUE\_EMPTY – Error, queue full

## Interrupt Management

### Nested interrupts

```
configMAX_SYSCALL_INTERRUPT_PRIORITY = 3  
configKERNEL_INTERRUPT_PRIORITY = 1
```



## Resources Management

### Sharing resources between tasks

To ensure data consistency is maintained at all times access to a resource that is shared between tasks, or between tasks and interrupts, must be managed using a 'mutual exclusion' technique.

**Critical sections:** regions of code

Start with **taskENTER\_CRITICAL()** and finish with **taskEXIT\_CRITICAL()**

```
/* Ensure access to the PORTA register cannot be interrupted by placing it within  
a critical section. Enter the critical section. */
```

```
taskENTER_CRITICAL();
```

```
PORTA |= 0x01;
```

```
/* Access to PORTA has finished, so it is safe to exit the critical section. */
```

```
taskEXIT_CRITICAL();
```

### Disable interrupts

## Resources Management

### Sharing resources between tasks

To ensure data consistency is maintained at all times access to a resource that is shared between tasks, or between tasks and interrupts, must be managed using a 'mutual exclusion' technique.

### Suspending the scheduler

#### Interrupts remain enabled

```
void vTaskSuspendAll( void );
```

The scheduler is suspended by calling `vTaskSuspendAll()`.  
Suspending the scheduler prevents a context switch from occurring

The scheduler is resumed (un-suspended) by calling `xTaskResumeAll()`.

```
BaseType_t xTaskResumeAll( void );
```



## Resources Management

### Sharing resources between tasks

To ensure data consistency is maintained at all times access to a resource that is shared between tasks, or between tasks and interrupts, must be managed using a 'mutual exclusion' technique.

### Suspending the scheduler

```
void vPrintString( const char *pcString )
{
    /* Write the string to stdout, suspending the scheduler as a method of mutual
    exclusion. */
    vTaskSuspendScheduler();
    {
        printf( "%s", pcString );
        fflush( stdout );
    }
    xTaskResumeScheduler();
}
```

## Resources Management

### Sharing resources between tasks

To ensure data consistency is maintained at all times access to a resource that is shared between tasks, or between tasks and interrupts, must be managed using a 'mutual exclusion' technique.

### Mutex

```
SemaphoreHandle_t xSemaphoreCreateMutex( void );
```

Returned value	If NULL is returned then the mutex could not be created because there is insufficient heap memory available.
----------------	--

## Resources Management

### Sharing resources between tasks

To ensure data consistency is maintained at all times access to a resource that is shared between tasks, or between tasks and interrupts, must be managed using a 'mutual exclusion' technique.

### Mutex

```
static void prvNewPrintString( const char *pcString )
{
    xSemaphoreTake( xMutex, portMAX_DELAY );
    {
        /* The following line will only execute once the mutex has been successfully
        obtained. */
        printf( "%s", pcString );
        fflush( stdout );
        /* The mutex MUST be given back! */
    }
    xSemaphoreGive( xMutex );
}
```

**Two risks: priority inversion or deadlock.**

## Resources Management

### Sharing resources between tasks

To ensure data consistency is maintained at all times access to a resource that is shared between tasks, or between tasks and interrupts, must be managed using a 'mutual exclusion' technique.

**Gatekeeper tasks:** provide a clean method of implementing mutual exclusion without the risk of priority inversion or deadlock.

Only the gatekeeper task is allowed to access the resource directly.

Indirect use of resources from other tasks by using the services of the gatekeeper.

## Resources Management

### Sharing resources between tasks

To ensure data consistency is maintained at all times access to a resource that is shared between tasks, or between tasks and interrupts, must be managed using a 'mutual exclusion' technique.

### Gatekeeper tasks

```
static void prvStdioGatekeeperTask( void *pvParameters )
{
char *pcMessageToPrint;
/* This is the only task that is allowed to write to standard out. */
for( ;; )
{
    /* The function will only return when a message has been successful
    received. */
    xQueueReceive( xPrintQueue, &pcMessageToPrint, portMAX_DELAY );
    /* Output the received string. */
    printf( "%s", pcMessageToPrint );
    fflush( stdout );
    /* Loop back to wait for the next message. */
}
}
```

QUESTIONS?